

LAKIREDDY BALI REDDY COLLEGE OF ENGINEERING

(AUTONOMOUS)



Department of Computer Science & Engineering

20AD53 - Machine Learning Lab

Name of the Student:

Registered Number:

Branch & Section: &/Sec

Academic Year: 2022 - 23

LAKIREDDY BALI REDDY COLLEGE OF ENGINEERING

(AUTONOMOUS)



CERTIFICATE

This is to certify that this is a bonafide record of the practical work
done by **Mr./Ms.**,
bearing Regd. Num.: 20761A05..... of B.Tech Semester, Branch,
Section in the **20AD53 - Machine Learning Lab** during the
Academic Year: 2022 - 2023 _____
No. of Experiments/Modules held: 12
No. of Experiments Done: 12

Date: / / 2022

Signature of the Faculty

INTERNAL EXAMINER

EXTERNAL EXAMINER

INDEX

S.No	DATE	EXPERIMENT	PAGE No.	REMARKS
1		Basic statistical functions for data exploration	4-5	
2.		Data Visualization: Box plot, scatter plot, histogram	5-6	
3.		Data Pre-processing: Handling missing values, outliers, normalization, Scaling	6-7	
4.		Principal Component Analysis (PCA)	8	
5.		Singular Value Decomposition (SVD)	9	
6.		Linear Discriminant Analysis (LDA)	10	
7.		Regression Analysis	11-13	
8.		Regularized Regression	14-15	
9.		K-Nearest Neighbour (KNN) Classifier	16-17	
10.		Support Vector Machines (SVMs)	18	
11.		Random Forest model	19	
12.		AdaBoost Classifier & XGBoost	20-22	

1. Basic statistical functions for data exploration

Program:

```
import pandas as pd

import numpy as np

from scipy import stats

from sklearn.datasets import load_boston

i=load_boston()

df=pd.DataFrame(i.data,columns=i.feature_names)

print(df)

print(df.describe())

# median & mode

median=[]

m=[]

for i in df.columns:

    median.append(np.median(df[i]))

    m.append(stats.mode(df[i]))

print(median,'\n',m)
```

OUTPUT:

```
The Boston housing prices dataset has an ethical problem. You can refer to
the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this
dataset unless the purpose of the code is to study and educate about
ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original
source::

import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]

Alternative datasets include the California housing dataset (i.e.
:func:`sklearn.datasets.fetch_california_housing`) and the Ames housing
dataset. You can load the datasets as follows::

from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()

for the California housing dataset and::

from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)

for the Ames housing dataset.
warnings.warn(msg, category=FutureWarning)
  CRIM      ZN  INDUS  CHAS    NOX     RAD   TAX    PTRATIO      B  LSTAT
0  0.00632  18.0    2.31   0.0  0.538  ...   1.0  296.0    15.3  396.90   4.98
1  0.02731   0.0    7.07   0.0  0.469  ...   2.0  242.0    17.8  396.90   9.14
2  0.02729   0.0    7.07   0.0  0.469  ...   2.0  242.0    17.8  392.83   4.03
3  0.03237   0.0    2.18   0.0  0.458  ...   3.0  222.0    18.7  394.63   2.94
4  0.06905   0.0    2.18   0.0  0.458  ...   3.0  222.0    18.7  396.90   5.33
```

```

Alternative datasets include the California housing dataset (i.e.
:func:`sklearn.datasets.fetch_california_housing`) and the Ames housing
dataset. You can load the datasets as follows::

from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()

for the California housing dataset and::

from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)

for the Ames housing dataset.
warnings.warn(msg, category=FutureWarning)

```

	CRIM	ZN	INDUS	CHAS	NOX	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	...	1.0	286.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	...	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	...	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	...	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	...	3.0	222.0	18.7	396.90
...
501	0.06263	0.0	11.93	0.0	0.573	...	1.0	273.0	21.0	391.99
502	0.04527	0.0	11.93	0.0	0.573	...	1.0	273.0	21.0	396.90
503	0.06076	0.0	11.93	0.0	0.573	...	1.0	273.0	21.0	396.90
504	0.10959	0.0	11.93	0.0	0.573	...	1.0	273.0	21.0	393.45
505	0.04741	0.0	11.93	0.0	0.573	...	1.0	273.0	21.0	396.90

```

[506 rows x 13 columns]

```

	CRIM	ZN	INDUS	...	PTRATIO	B	LSTAT
count	506.000000	506.000000	506.000000	...	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	...	18.455534	356.674032	12.653063
std	8.601545	23.322453	6.860353	...	2.164946	91.294864	7.141062
min	0.006320	0.000000	0.460000	...	12.600000	0.320000	1.730000
25%	0.022045	0.000000	5.190000	...	17.400000	375.377500	6.950000
50%	0.256510	0.000000	9.690000	...	19.050000	391.440000	11.360000
75%	3.677083	12.500000	18.100000	...	20.200000	396.225000	16.955000
max	88.976200	100.000000	27.740000	...	22.000000	396.900000	37.970000

```

[8 rows x 13 columns]
[0.25651, 0.0, 9.69, 0.0, 0.538, 6.2085, 77.5, 3.2074499999999997, 5.0, 330.0, 19.05, 391.44, 11.36]
[ModeResult(mode=array([0.01501]), count=array([2])), ModeResult(mode=array([18.1]), count=array([132])), ModeResult(mode=array([0.])
, count=array([471])), ModeResult(mode=array([0.538]), count=array([23])), ModeResult(mode=array([5.713]), count=array([3])), ModeResult(mode=array([100.]), count=array([43])), ModeR
esult(mode=array([3.4952]), count=array([5])), ModeResult(mode=array([24.]), count=array([132])), ModeResult(mode=array([666.]), count=array([132])), ModeResult(mode=array([20.2]), co
unt=array([140])), ModeResult(mode=array([396.9]), count=array([121])), ModeResult(mode=array([6.36]), count=array([3]))]

```

2. Data Visualization: Box plot, scatter plot, histogram

Program:

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import load_boston

i=load_boston()

df=pd.DataFrame(i.data,columns=i.feature_names)

df['CRIM'].plot(kind='box')

plt.show()

plt.scatter(df['CRIM'],df['ZN'],cmap='viridis',c=df['LSTAT'])


plt.show()

df['LSTAT'].hist().grid(False)

plt.show()

```

OUTPUT:



3. Data Pre-processing: Handling missing values, outliers, normalization, Scaling

Program:

```
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
import pandas as pd
import numpy as np
from scipy import stats
from sklearn.datasets import load_boston
i=load_boston()
df=pd.DataFrame(i.data,columns=i.feature_names)
print(df.isna().sum(),'\n')
#checking outliers
q1=df['ZN'].quantile(0.25)
```

```

q3=df['ZN'].quantile(0.75)
IQR=q3-q1
Upper=q3+1.5*IQR
Lower=q1-1.5*IQR
df[df['ZN'] > Upper]
df[df['ZN'] < Lower]
new_df = df[df['ZN'] > Upper ]
plt.subplot(1,2,1)
sns.boxplot(x=df['ZN'])
plt.subplot(1,2,2)
sns.boxplot(x=new_df['ZN'])
#normalize the data
s=StandardScaler()
df=s.fit_transform(df)
df


```

OUTPUT:

```

Out[2]: array([[ -0.41978194,  0.28482986, -1.2879095 , ..., -1.45900038,
         0.44105193, -1.0755623 ],
        [ -0.41733926, -0.48772236, -0.59338101, ..., -0.30309415,
         0.44105193, -0.49243937],
        [ -0.41734159, -0.48772236, -0.59338101, ..., -0.30309415,
         0.39642699, -1.2087274 ],
        ...,
        [ -0.41344658, -0.48772236,  0.11573841, ...,  1.17646583,
         0.44105193, -0.98304761],
        [ -0.40776407, -0.48772236,  0.11573841, ...,  1.17646583,
         0.4032249 , -0.86530163],
        [ -0.41500016, -0.48772236,  0.11573841, ...,  1.17646583,
         0.44105193, -0.66905833]])

```




4. Principal Component Analysis (PCA)

Program:

```
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.datasets import load_iris
iris=load_iris()
x,y=iris.data,iris.target
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.fit_transform(x_test)
x_1=x_train
pca=PCA(n_components=2)
x_train=pca.fit_transform(x_train)
comps=pca.components_
print("components:\n",pca.components_)
print("Co variance\n",pca.get_covariance())
plt.figure(figsize=(10,6))
plt.subplot(1,3,1)
plt.plot(x_1)
plt.subplot(1,3,2)
plt.plot(x_train)
plt.subplot(1,3,3)
c=["black",'yellow','red','green']
for i in range(0,len(comps)):
    for j in range(0,len(comps[i])):
        plt.scatter(i,comps[i][j],c=c[i])
```

OUTPUT:

```
components:
[[ 0.53412301 -0.22791604  0.58442075  0.5667621 ]
 [ 0.30791713  0.94856275  0.01976354  0.07088833]]
Co variance
[[ 0.97606998 -0.09317132  0.88720786  0.87407504]
 [-0.09317132  1.00793386 -0.36025464 -0.30724325]
 [ 0.88720786 -0.36025464  1.05395785  0.93708159]
 [ 0.87407504 -0.30724325  0.93708159  1.00049985]]
```



5. Singular Value Decomposition (SVD)

Program:

```
from numpy.linalg import svd
A = np.array([[3,4,3],[1,2,3],[4,2,1]])
S,U,VT=svd(A)
print("Singular Matrix:\n",S,\n',"U:\n",U,\n',"Vector Matrix:\n",VT)
```

OUTPUT:

Singular Matrix:

```
[[-0.73553325 -0.18392937 -0.65204358]
 [-0.42657919 -0.62196982  0.65664582]
 [-0.52632788  0.76113306  0.37901904]]
```

U:

```
[7.87764972 2.54031671 0.69958986]
```

Vector Matrix:

```
[[-0.60151068 -0.61540527 -0.5093734 ]
 [ 0.73643349 -0.18005275 -0.65210944]
 [ 0.30959751 -0.76737042  0.5615087  ]]
```

6. Linear Discriminant Analysis (LDA)

Program:


```
from sklearn.datasets import load_iris
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

data=load_iris()
x,y=data.data,data.target
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)
ss=StandardScaler()
x_train=ss.fit_transform(x_train)
x_test=ss.fit_transform(x_test)
lda=LDA(n_components=2)
X_train=lda.fit_transform(x_train,y_train)
X_test=lda.transform(x_test)
plt.figure(figsize=(7,7))

plt.subplot(2,2,1)
plt.plot(x_train)
plt.subplot(2,2,2)
plt.plot(X_train)
plt.subplot(2,2,3)
plt.plot(x_test)
plt.subplot(2,2,4)
plt.plot(X_test)
```

OUTPUT:

```
Out[5]: [<matplotlib.lines.Line2D at 0x7f8cb53cfa0>,
<matplotlib.lines.Line2D at 0x7f8cb53d9040>]
```



7. Regression Analysis:

7 (A) .Linear Regression

Program:

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

import seaborn as sns

from sklearn.datasets import make_blobs
x,y=make_blobs(n_samples=40,n_features=1)

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)

lr=LinearRegression()

lr.fit(x_train,y_train) print("slope:",lr.coef_)

print("Intercept:",lr.intercept_)

print("Score:",lr.score(x_train,y_train))

y_pred=lr.predict(x_test)

print('Accuracy:',r2_score(y_test,y_pred))

plt.scatter(x_test,y_test)

plt.plot(x_test,y_pred,c='green',marker='o',markerfacecolor='red')

plt.xlabel("X")


plt.ylabel("Y")
```

OUTPUT:

```
plt.plot(x_test,y_pred,c='green',marker='o',markerfacecolor='red')
plt.xlabel("X")
plt.ylabel("Y")
```

```
slope: [0.00814197]
Intercept: 0.9789237089972992
Score: 0.00012075418594370557
Accuracy: -7.164353685640279e-05
```

Out[6]: Text(0, 0.5, 'Y')



7 (B) .Logistic Regression

Program:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
import seaborn as sns

df = sns.load_dataset('iris')
x = df[['sepal_length', 'sepal_width', 'petal_length', 'petal_width']]
y = df['species']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=1)
lrc = LogisticRegression(solver='liblinear', random_state=0)
lrc.fit(x_train, y_train)


y_pred = lrc.predict(x_test)
print("classes:", lrc.classes_)
print("coefficient:\n", lrc.coef_)
print("Intercept:", lrc.intercept_)
print("Score:", lrc.score(x_train, y_train))
print("Predicted Probabilities:", lrc.predict_proba(x_test[:1]))
print("accuracy:", accuracy_score(y_test, y_pred))

sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, xticklabels=pd.unique(y_train), yticklabels=pd.unique(y_train))
```

OUTPUT:

```
classes: ['setosa' 'versicolor' 'virginica']
coefficient:
[[ 0.40624466  1.34920549 -2.12162959 -0.94996585]
 [ 0.628089   -1.72885637  0.30292395 -0.99777747]
 [-1.6655058  -0.95801685  2.19808191  2.08216345]]
Intercept: [ 0.26653163  0.67420401 -0.86971453]
Score: 0.9523809523809523
Predicted Probabilities: [[9.25059393e-01 7.49291898e-02 1.14167043e-05]]
accuracy: 0.8888888888888888
```

Out[7]: <AxesSubplot:>



7 C). Polynomial Regression

Program:

```
import pandas as pd
import numpy as np

from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression

df=pd.read_csv('data.csv')
x=df[['Level']]
y=df['Salary']

poly=PolynomialFeatures(degree=2)
x_poly=poly.fit_transform(x)
lin=LinearRegression().fit(x,y)
y_pred=lin.predict(x)

lin2=LinearRegression().fit(x_poly,y)
y_pred1=lin2.predict(x_poly)

plt.figure(figsize=(7,5)) plt.subplot(1,2,1)
plt.scatter(x,y)


plt.plot(x,y_pred,marker='o',markerfacecolor='yellow',c='red')
plt.subplot(1,2,2)

plt.scatter(x,y)

plt.plot(x,y_pred1,marker='o',markerfacecolor='yellow',c='red')
```

OUTPUT:

[<matplotlib.lines.Line2D at 0x7f8caffff1730>]



8. Regularized Regression

Program:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler,LabelEncoder,scale
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge,Lasso
from sklearn.metrics import r2_score
import seaborn as sns
df=sns.load_dataset('flights')
le=LabelEncoder()
df['month']=le.fit_transform(df['month'])
x=df[['year','month']]
y=df['passengers']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)
ridge=Ridge(alpha=0.3)
ridge.fit(x_train,y_train)
y_pred=ridge.predict(x_test)
lasso=Lasso(alpha=0.1)
lasso.fit(x_train,y_train)
y_pred1=lasso.predict(x_test)
print('Ridge Regularized Regression:')
print('slope:',ridge.coef_)
print('Intercept:',ridge.intercept_)
print('score:',ridge.score(x_train,y_train))
print('accuracy:',r2_score(y_test,y_pred),'\n')
print('Lasso Regularized Regression:')
print('slope:',lasso.coef_)
print('Intercept:',lasso.intercept_)
print('score:',lasso.score(x_train,y_train))
print('accuracy:',r2_score(y_test,y_pred1))
plt.subplot(1,2,1)
```

```

plt.title('Ridge Regularized')
plt.scatter(x_test['month'][:10],y_test[:10])
plt.plot(x_test['month'][:10],y_pred[:10],c='red',marker='o',markerfacecolor='yellow')
plt.subplot(1,2,2)
plt.title('Lasso Regularized')
plt.scatter(x_test['month'][:10],y_test[:10])
plt.plot(x_test['month'][:10],y_pred1[:10],c='red',marker='o',markerfacecolor='yellow')

```

OUTPUT:

```

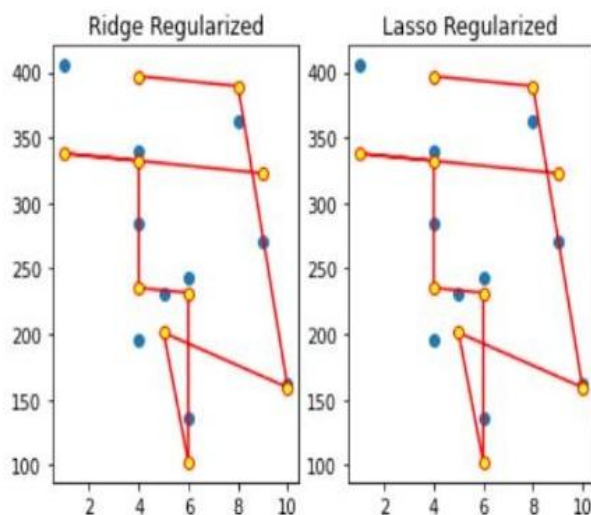
plt.scatter(x_test['month'][:10],y_test[:10])
plt.plot(x_test['month'][:10],y_pred1[:10],c='red',marker='o',markerfacecolor='yellow')

```

Ridge Regularized Regression:
slope: [32.3137578 -1.90613899]
Intercept: -62865.90705172744
score: 0.8572612820122529
accuracy: 0.8239841277324573

Lasso Regularized Regression:
slope: [32.31311528 -1.89857863]
Intercept: -62864.692308705315
score: 0.8572612218618013
accuracy: 0.8240219138855976

Out[10]: [<matplotlib.lines.Line2D at 0x7f8cafe9dd00>]



9. K-Nearest Neighbour (KNN) Classifier

Program:

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np
import seaborn as sns

df=sns.load_dataset('tips')
l=LabelEncoder()
df['sex']=l.fit_transform(df['sex'])
df['smoker']=l.fit_transform(df['smoker'])
df['day']=l.fit_transform(df['day'])
df['time']=l.fit_transform(df['time'])
x=df[['total_bill','tip','sex','smoker','day','time']]
y=df['size']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)
D1=[]
D2=[]
D3=[]
D4=[]
for i in range(1,11):
    knn=KNeighborsClassifier(n_neighbors=i)
    knn1=KNeighborsClassifier(p=1,n_neighbors=i)
    knn2=KNeighborsClassifier(metric='cosine',algorithm='brute',n_neighbors=i)
    knn3=KNeighborsClassifier(metric='jaccard',n_neighbors=i)
    knn.fit(x_train,y_train)
    knn1.fit(x_train,y_train)
    knn2.fit(x_train,y_train)
    knn3.fit(x_train,y_train)
    D1.append(accuracy_score(y_test,knn.predict(x_test)))
    D2.append(accuracy_score(y_test,knn1.predict(x_test)))
    D3.append(accuracy_score(y_test,knn2.predict(x_test)))
```



```

D4.append(accuracy_score(y_test,knn3.predict(x_test)))
D=np.array((D1,D2,D3,D4)).T
dff=pd.DataFrame(D,columns=['Manhattan','Euclidean','Cosine','Jaccard'],index=[i
for i in range(1,11)])
print(dff)
plt.subplot(1,2,1)
plt.title('size')
df['size'].value_counts().plot(kind='bar',cmap='rainbow')
plt.subplot(1,2,2)
plt.title('Accuracys')
np.max(dff).plot(kind='bar')


```

OUTPUT:

	Manhattan	Euclidean	Cosine	Jaccard
1	0.581081	0.594595	0.486486	0.527027
2	0.621622	0.608108	0.594595	0.527027
3	0.608108	0.581081	0.581081	0.621622
4	0.608108	0.594595	0.594595	0.621622
5	0.594595	0.608108	0.608108	0.621622
6	0.594595	0.635135	0.608108	0.621622
7	0.608108	0.635135	0.621622	0.621622
8	0.594595	0.621622	0.635135	0.621622
9	0.608108	0.635135	0.608108	0.621622
10	0.621622	0.662162	0.662162	0.621622

/home/shesiram/20761A0576/lib/python3.8/site-packages/numpy/core/fromnumeric.py:84: FutureWarning: In a future version, DataFrame.max(axis=None) will return a scalar max over the entire DataFrame. To retain the old behavior, use 'frame.max(axis=0)' or just 'frame.max()'
return reduction(axis=axis, out=out, **passkwargs)

Out[12]: <AxesSubplot:title={'center':'Accuracys'}>

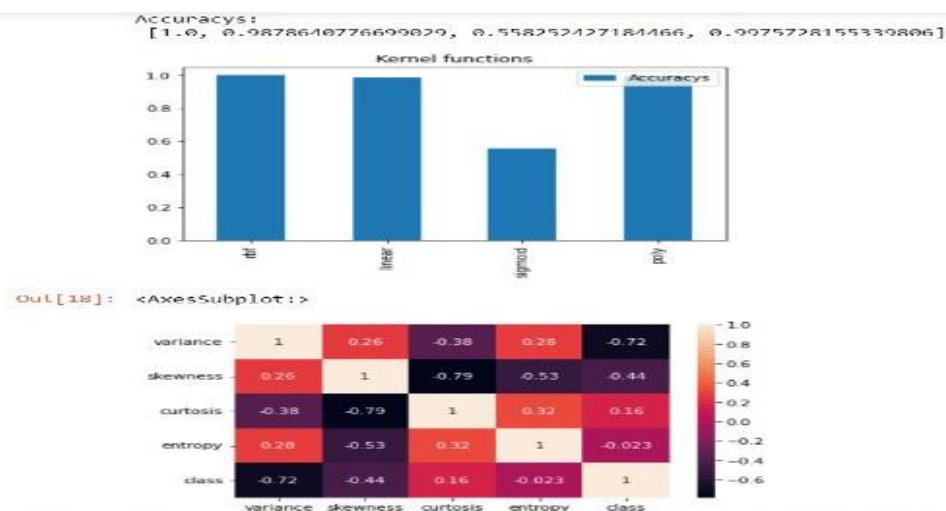


10. Support Vector Machines (SVMs)

Program:

```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
df=pd.read_csv('BankNote_Authentication.csv')
x=df[['variance', 'skewness', 'curtosis', 'entropy']]
y=df['class']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
s=[]
kernel=['rbf','linear','sigmoid','poly']
for i in kernel:
    svm=SVC(kernel=i,gamma='auto',C=1)
    svm.fit(x_train,y_train)
    s.append(accuracy_score(y_test,svm.predict(x_test)))
print("Accuracys:\n",s)
df1=pd.DataFrame(np.array(s).reshape(4,1),columns=['Accuracys'])
df1.plot(kind='bar')
plt.title('Kernel functions')
plt.xticks(range(4),kernel)
plt.show()
sns.heatmap(df.corr(),annot=True)
```

OUTPUT:



11. Random Forest model

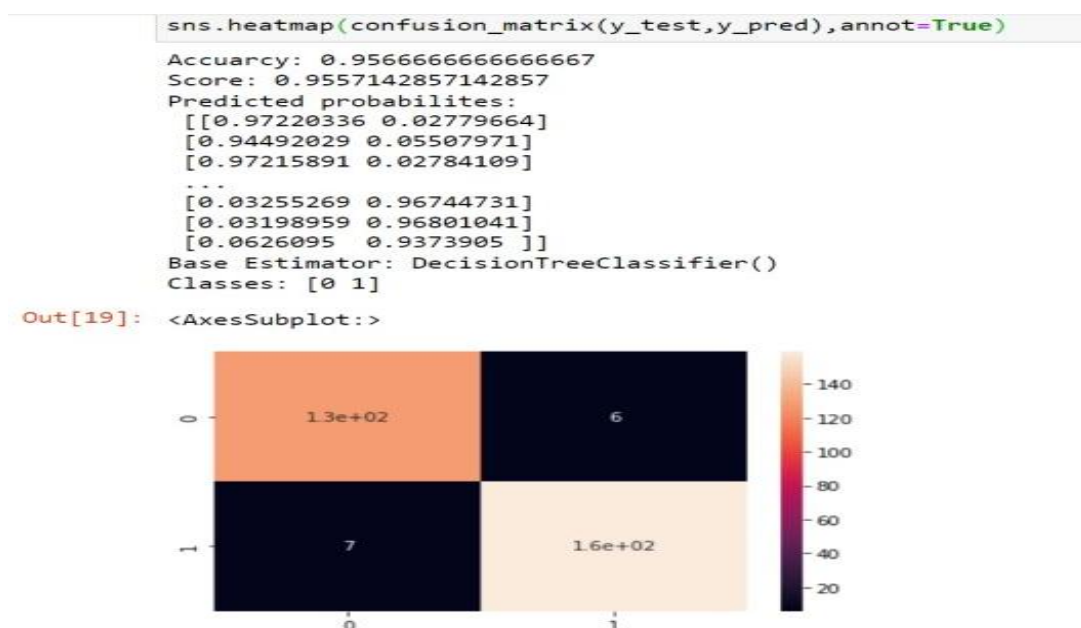
Program:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.datasets import make_classification
import seaborn as sns
from sklearn.metrics import accuracy_score, confusion_matrix

x,y=make_classification(n_samples=1000,n_features=4,random_state=0,shuffle=False)

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
rfc=RandomForestClassifier(max_depth=2,n_estimators=1000,random_state=0)
rfc.fit(x_train,y_train)
y_pred=rfc.predict(x_test)
print('Accuracy:',accuracy_score(y_test,y_pred))
print('Score:',rfc.score(x_train,y_train))
print('Predicted probabilities:\n',rfc.predict_proba(x))
print('Base Estimator:',rfc.base_estimator_)
print('Classes:',rfc.classes_)
sns.heatmap(confusion_matrix(y_test,y_pred),annot=True)
```

OUTPUT:



12.(a)AdaBoost Classifier

Program:

```
from sklearn.ensemble import AdaBoostClassifier,RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import seaborn as sns
df=sns.load_dataset('iris')
df['species']=LabelEncoder().fit_transform(df['species'])
x=df[['sepal_length', 'sepal_width', 'petal_length', 'petal_width']]
y=df['species']
svm=SVC()
lrc=LogisticRegression()
dtc=DecisionTreeClassifier()
rfc=RandomForestClassifier()
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
classifier=['LogisticRegression','DecisionTreeClassifier','SVC','RandomForestClassifier']
instance=[lrc,dtc,svm,rfc]
s=[]
for i in instance:
    ada=AdaBoostClassifier(n_estimators=100,base_estimator=i,learning_rate=1.0,random_state=0,algorithm='SAMME')
    ada.fit(x_train,y_train)
    s.append(accuracy_score(y_test,ada.predict(x_test)))
```

```

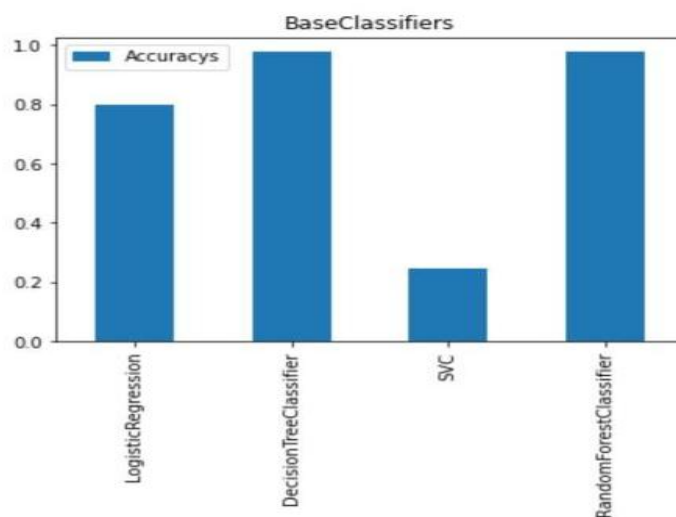
print("Accuracys:\n",s)
df1=pd.DataFrame(np.array(s).reshape(4,1),columns=['Accuracys'])
df1.plot(kind='bar')
plt.title('BaseClassifiers')
plt.xticks(range(4),classifier)
plt.show()
sns.heatmap(df.corr(),annot=True,cmap='rainbow')

```


OUTPUT:

```
sns.heatmap(df.corr(),annot=True,cmap='rainbow')
```

Accuracys:
[0.8, 0.9777777777777777, 0.24444444444444444, 0.9777777777777777]



Out[22]: <AxesSubplot:>



(b). XGBoost

Program:

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import seaborn as sns
from sklearn.metrics import accuracy_score, confusion_matrix
from xgboost import XGBClassifier as XGB
from sklearn.preprocessing import LabelEncoder
df=sns.load_dataset('iris')
le=LabelEncoder()
df['species']=le.fit_transform(df['species'])
x=df[['sepal_length', 'sepal_width', 'petal_length', 'petal_width']]
y=df['species']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
xgb=XGB()
xgb.fit(x_train,y_train)
y_pred=xgb.predict(x_test)
sns.heatmap(df.corr(),annot=True,cmap='viridis')
print('Accuracy:',accuracy_score(y_test,y_pred))
print('Predicted values:\n',y_pred)
```

OUTPUT:

